

DataStax

Developers

Apache Cassandra™ Data Modelling

Building efficient applications with Apache Cassandra





Developer Advocate Lead

dtsx.io/aleks



@aleks-volochnev
@HadesArchitect

- IT Exorcist
- Apache Cassandra™ certified
- Cloud Architect certified



Aleksandr Volochnev



Developer Advocate



@ArtemChebotko

- Data professional, computer scientist
- Data modeling, data quality, data warehousing, data analytics
- Author of the Cassandra Data Modeling Methodology
- Google Cloud Certified Data Engineer

Astra DB

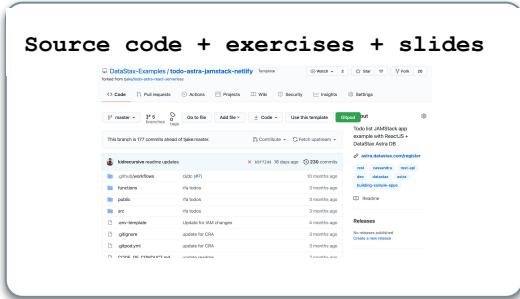
Astra
STREAMING

PULSAR



Artem Chebotko

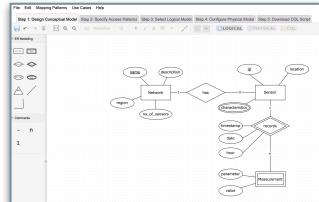
Nothing to install !



SCENARIOS

K L L R
C O D A BETA

DATA MODELING: KDM



<http://kdm.kashliev.com>

Hands-On Housekeeping

01



Data Storage Overview

Data Distribution & Organization

02

Key Definition

Primary, Partition, Clustering

03

CQL Data Types

Brief Review

04

Data Modeling Process

Methodology, Notation, Tool

05

Optimization Techniques

Five important optimizations

06

What's next?

Homework, Next Session



Agenda

01



Data Storage Overview
Data Distribution & Organization

02

Key Definition
Primary, Partition, Clustering

03

CQL Data Types
Brief Review

04

Data Modeling Process
Methodology, Notation, Tool

05

Optimization Techniques
Five important optimizations

06

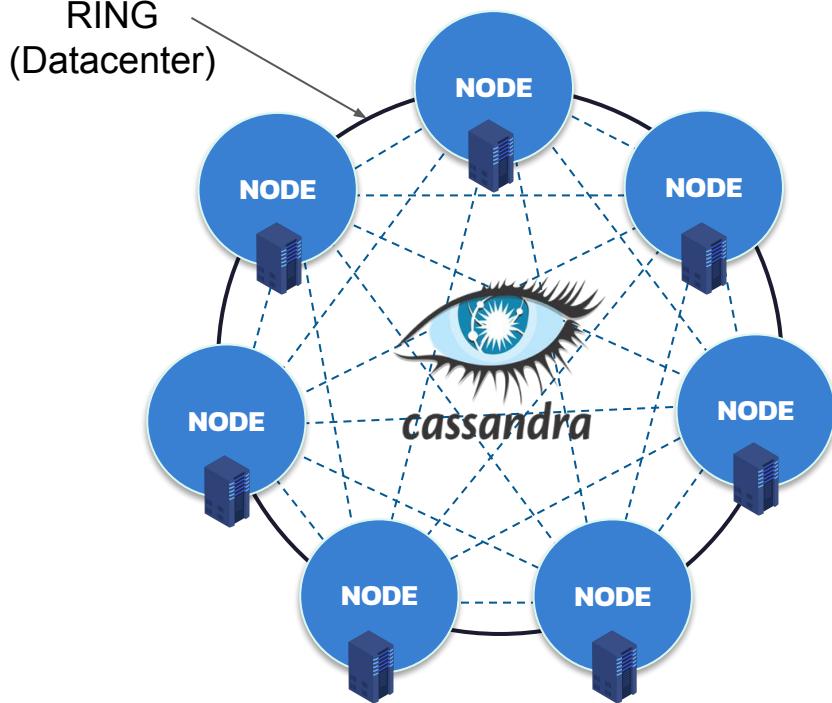
What's next?
Homework, Next Session



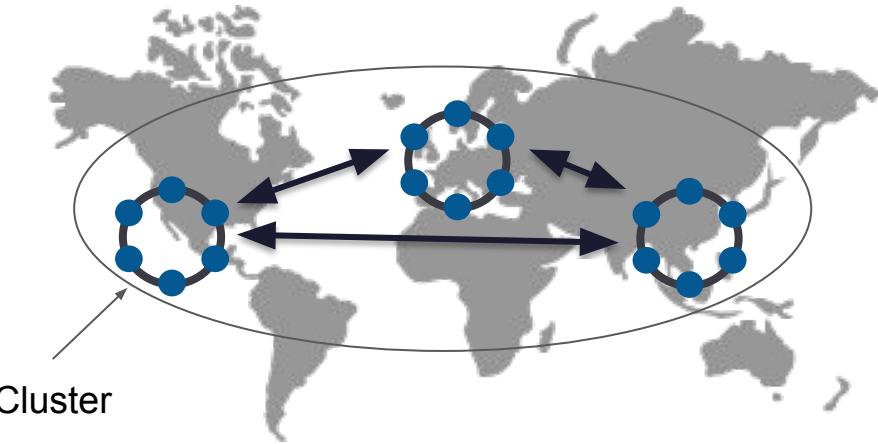
Agenda

Data Distribution

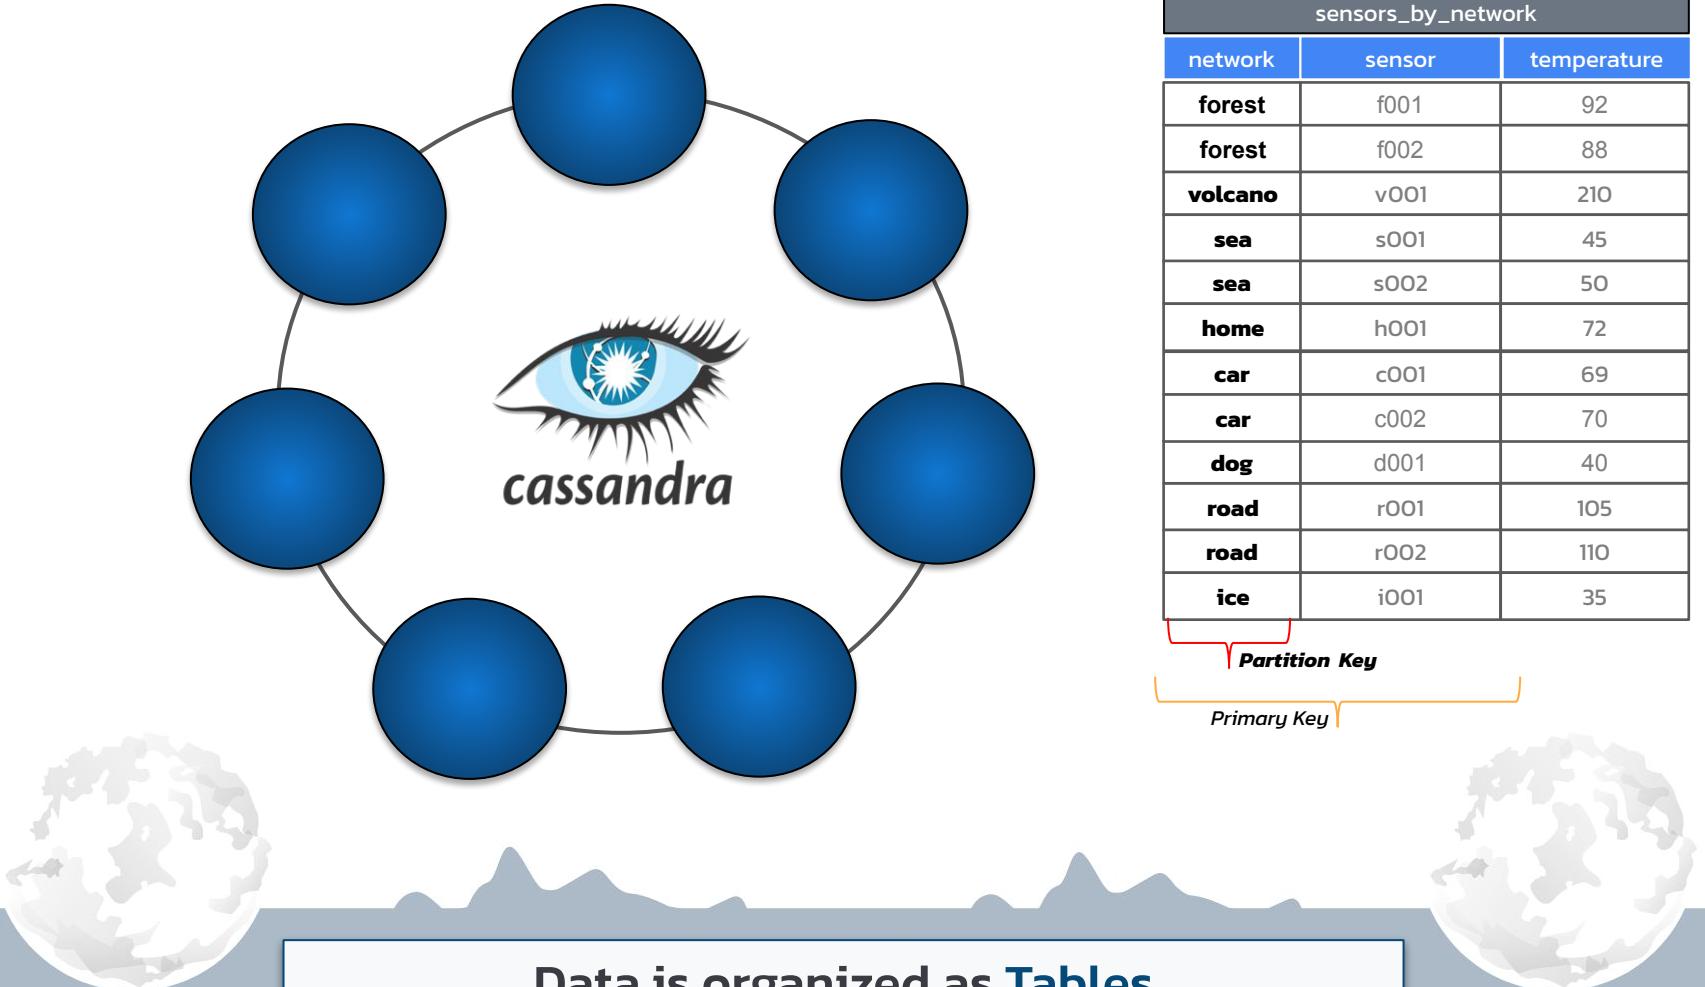


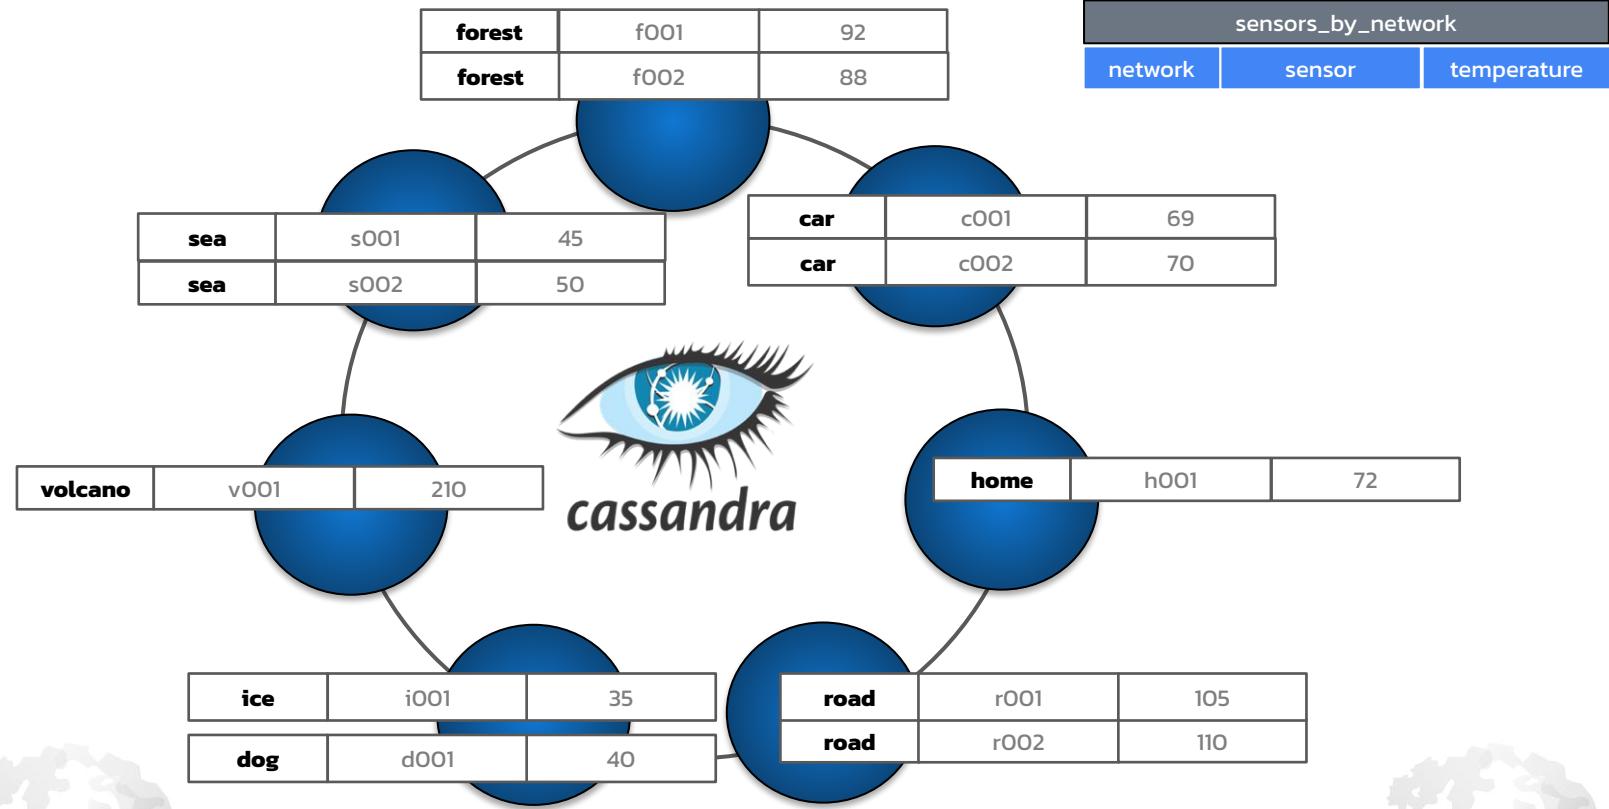


1. **NO** Single Point of Failure (masterless)
2. Scales for writes and reads
3. Application can contact any node
(in case of failure - just contact next one)

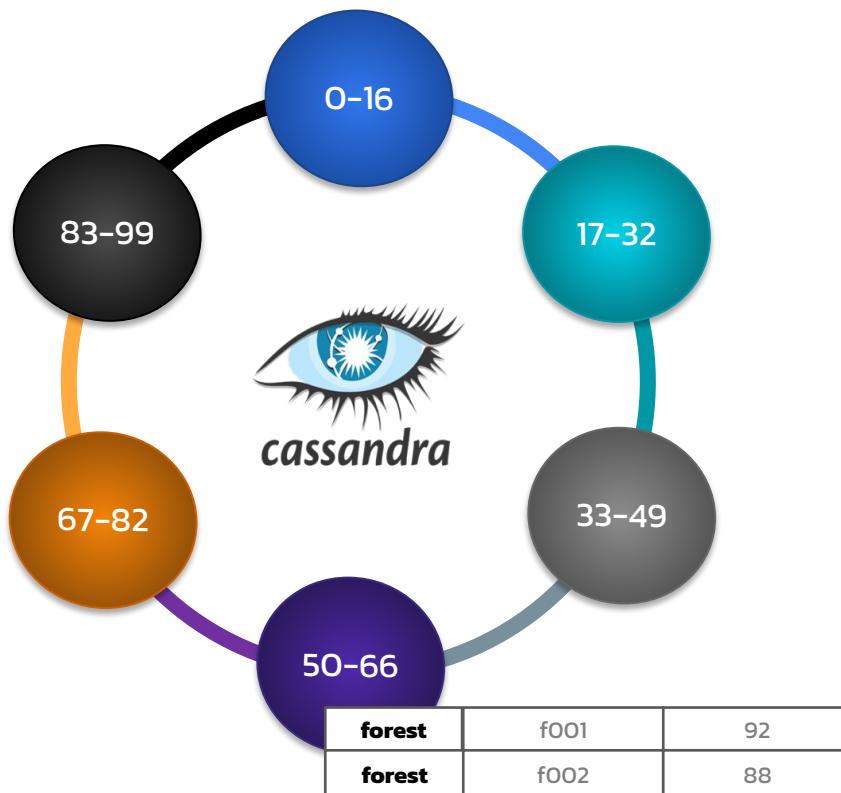
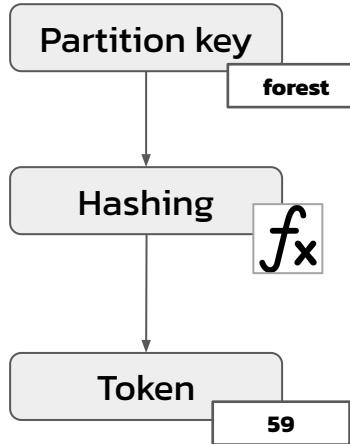


Apache Cassandra™ Infrastructure





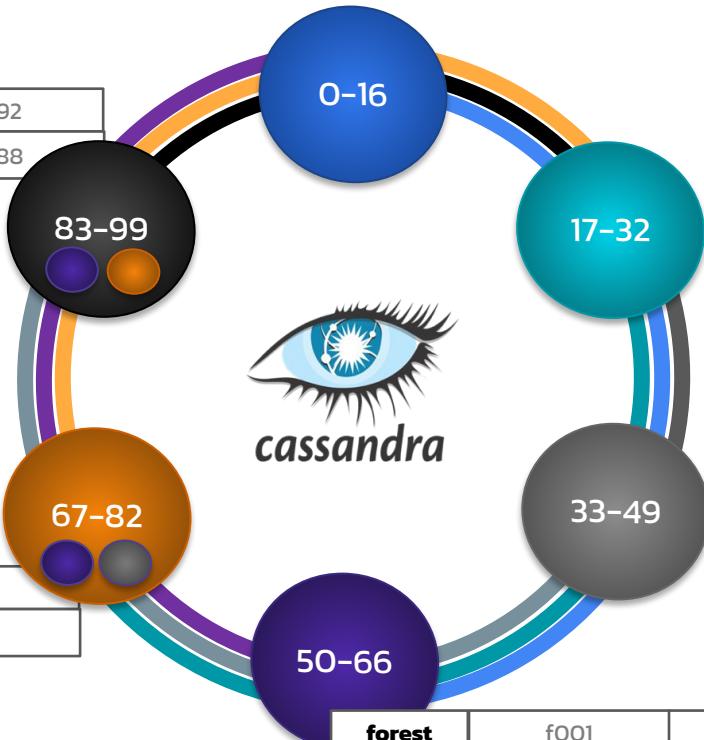
Data is organized as Distributed tables



Per Table, each node owns a range of tokens

RF = 3

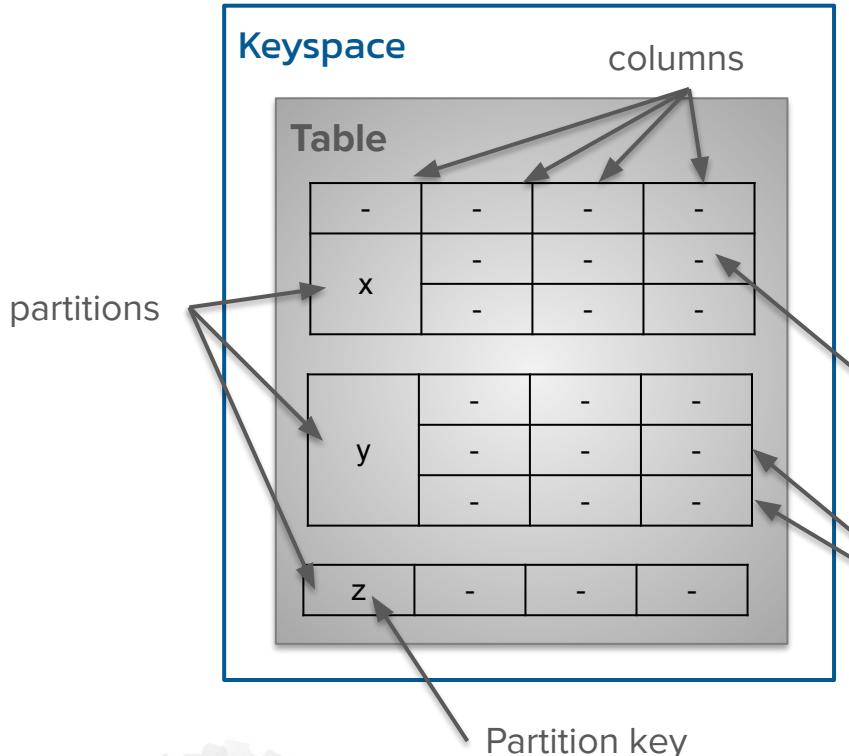
forest	f001	92
forest	f002	88



Data is replicated X times (replication factor)

Data Organization





- Tabular data model, with one twist
- *Keyspaces contain tables*
- *Tables are organized in rows and columns*
- Groups of related rows called *partitions* are stored together on the same node (or nodes)
- Each row contains a *partition key*
 - One or more columns that are hashed to determine which node(s) store that data

Cell \in Row \in Partition \in Table \in Keyspace

```
CREATE KEYSPACE sensor_data  
WITH REPLICATION = {  
    'class' : 'NetworkTopologyStrategy',  
    'us-west-1' : 3,  
    'eu-central-1' : 5  
};
```

keyspace

replication strategy

Replication factor by data center



Create a Keyspace

```
Keyspace  
↓  
CREATE TABLE sensor_data.sensors_by_network (  
    network      text,  
    sensor       text,  
    temperature integer,  
    PRIMARY KEY ((network), sensor)  
);
```

Table
↓

Primary key

Partition key

Clustering columns



Create a table



Lab 1

Create Your Astra DB Instance

[https://github.com/datastaxdevs/
workshop-cassandra-data-modeling](https://github.com/datastaxdevs/workshop-cassandra-data-modeling)

01



Data Storage Overview

Data Distribution & Organization

02

Key Definition

Primary, Partition, Clustering

03

CQL Data Types

Brief Review

04

Data Modeling Process

Methodology, Notation, Tool

05

Optimization Techniques

Five important optimizations

06

What's next?

Homework, Next Session



Agenda

Key Definition



An identifier for a partition.
Consists of at least one column,
may have more if needed.

PARTITIONS TABLE.

Defines data distribution strategy

```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
);
```

Partition key

Examples:

```
PRIMARY KEY ((sensor), timestamp);
```

```
PRIMARY KEY ((sensor), date, timestamp);
```

```
PRIMARY KEY ((sensor, timestamp));
```

Partition Key

Used to **ensure uniqueness** and **sorting order**. Optional.

Defines physical ordering on disk

```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```

Clustering columns

PRIMARY KEY ((**sensor**));

Not Unique

PRIMARY KEY ((**sensor**), timestamp);

Not filter on dates

PRIMARY KEY ((**sensor**), value, timestamp);

Not sorted

PRIMARY KEY ((**sensor**), date, timestamp);



Clustering Columns

An identifier for a row. Consists of at least one Partition Key and zero or more Clustering Columns.

UNIQUELY IDENTIFIES A ROW

```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```

Primary key

Examples:

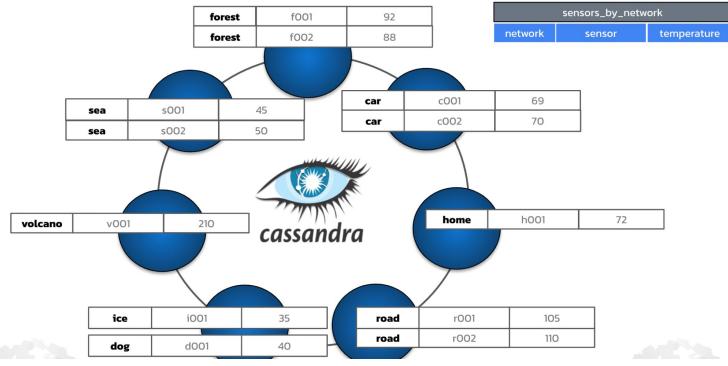
```
PRIMARY KEY ((sensor), timestamp);
```

```
PRIMARY KEY ((sensor, date), timestamp);
```

Primary Key

Partition Keys:

Defines data distribution over the cluster



Clustering Columns

Defines how data is physically stored (written on disk)

```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```

```
PRIMARY KEY ((sensor, date), timestamp);
```

```
SELECT * FROM temperatures_by_sensor ...
```

```
WHERE sensor = ?;
```

```
WHERE sensor > ?;
```

```
WHERE sensor = ? AND date > ?:
```

```
WHERE sensor = ? AND date = ? AND timestamp > ?;
```

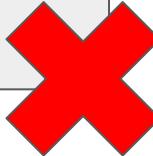
Important:

Once created, the data model cannot be changed! You will need new tables and migration. Stay lazy, design it right in advance!

```
ALTER TABLE temperatures_by_sensor  
ADD season TEXT;
```



```
ALTER TABLE temperatures_by_sensor  
DROP PRIMARY KEY  
ADD PRIMARY KEY (sensor, timestamp)
```



Consequence II: Schema Immutability

QUESTION

Why don't we use sequential IDs in Cassandra?



[My]SQL

```
CREATE TABLE table_name (
    column1 integer NOT NULL AUTO_INCREMENT,
    ...
);
```

```
mysql> INSERT INTO table_name ...
mysql> SELECT LAST_INSERT_ID();
```

Result: 99 (last value + 1)

Usually ID generated by a database

Cassandra

```
CREATE TABLE todoitems (
    item_id UUID,
    ...
);
```

```
cql> INSERT INTO todoitems ( item_id )
      VALUES ( UUID() );
cql> SELECT item_id FROM todoitems ...
```

Result: 123e4567-e89b-12d3-a456-426614174000

Usually UUID generated in an application code

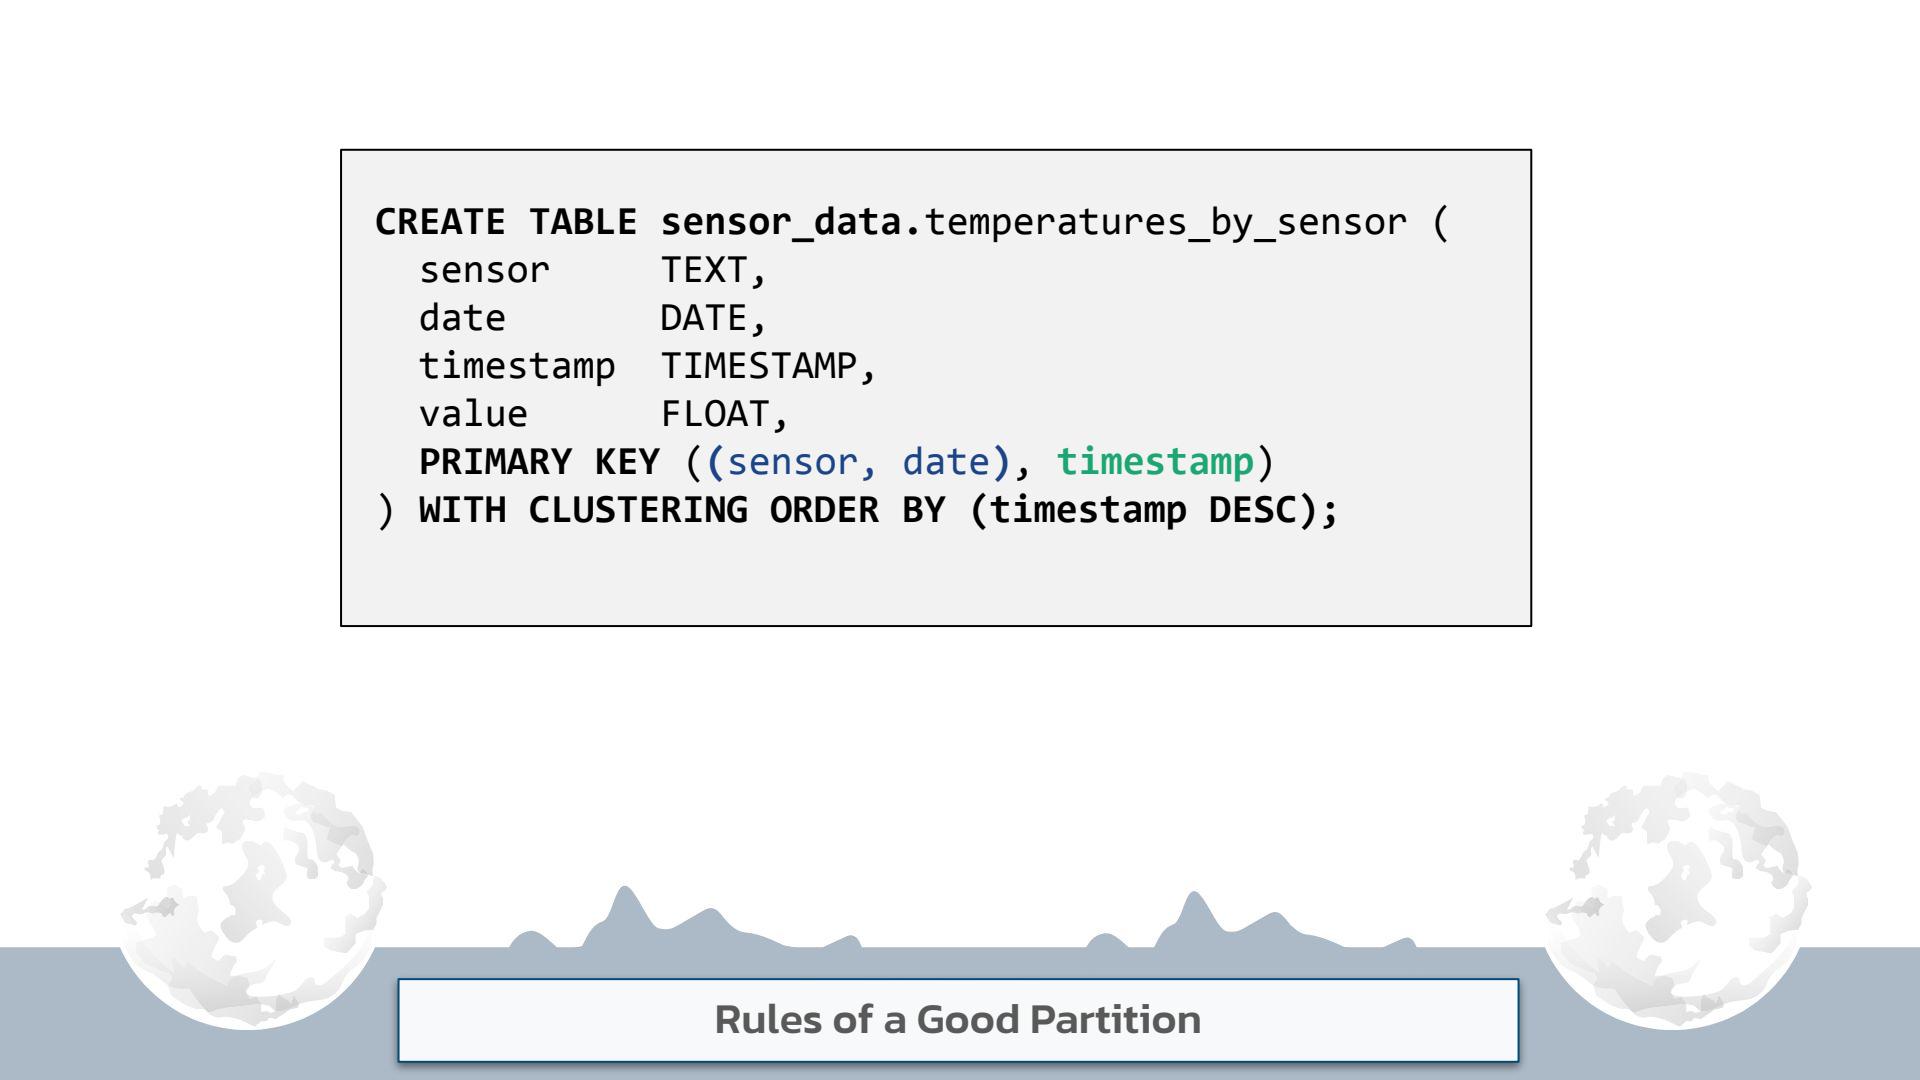


Different Approaches

Rules for good partitioning



```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```



Rules of a Good Partition

- ❖ **Store together what you retrieve together**
- ❖ Avoid big partitions
- ❖ Avoid hot partitions

Q: Show temperature evolution over time for **sensor X** On july 6th 2022

```
PRIMARY KEY ((sensor, timestamp));
```



```
PRIMARY KEY ((sensor), timestamp);
```



Rules of a Good Partition

- ❖ Store together what you retrieve together
- ❖ **Avoid big partitions**
- ❖ Avoid hot partitions

BUCKETING

PRIMARY KEY ((sensor), timestamp);



PRIMARY KEY ((sensor, month), timestamp);



- Up to 2 billion cells per partition
- Up to ~100k values in a partition
- Up to ~100MB in a Partition

Rules of a Good Partition

- ❖ Store together what you retrieve together
- ❖ Avoid big partitions
- ❖ **Avoid hot partitions**

```
PRIMARY KEY ((date), sensor, timestamp);
```



```
PRIMARY KEY ((date, sensor), timestamp);
```





Lab 2

Dynamic Bucketing

[https://github.com/datastaxdevs/
workshop-cassandra-data-modeling](https://github.com/datastaxdevs/workshop-cassandra-data-modeling)



01



Data Storage Overview

Data Distribution & Organization

02

Key Definition

Primary, Partition, Clustering

03

CQL Data Types

Brief Review

04

Data Modeling Process

Methodology, Notation, Tool

05

Optimization Techniques

Five important optimizations

06

What's next?

Homework, Next Session



Agenda

type	constants supported	description
ascii	<code>string</code>	ASCII character string
bigint	<code>integer</code>	64-bit signed long
blob	<code>blob</code>	Arbitrary bytes (no validation)
boolean	<code>boolean</code>	Either <code>true</code> or <code>false</code>
counter	<code>integer</code>	Counter column (64-bit signed value). See Counters for details
date	<code>integer</code> , <code>string</code>	A date (with no corresponding time value). See Working with dates below for details
decimal	<code>integer</code> , <code>float</code>	Variable-precision decimal
double	<code>integer</code> <code>float</code>	64-bit IEEE-754 floating point
duration	<code>duration</code> ,	A duration with nanosecond precision. See Working with durations below for details
float	<code>integer</code> , <code>float</code>	32-bit IEEE-754 floating point
inet	<code>string</code>	An IP address, either IPv4 (4 bytes long) or IPv6 (16 bytes long). Note that there is no <code>inet</code> constant, IP address should be input as strings
int	<code>integer</code>	32-bit signed int
smallint	<code>integer</code>	16-bit signed int
text	<code>string</code>	UTF8 encoded string
time	<code>integer</code> , <code>string</code>	A time (with no corresponding date value) with nanosecond precision. See Working with times below for details
timestamp	<code>integer</code> , <code>string</code>	A timestamp (date and time) with millisecond precision. See Working with timestamps below for details
timeuuid	<code>uuid</code>	Version 1 UUID , generally used as a “conflict-free” timestamp. Also see Timeuuid functions
tinyint	<code>integer</code>	8-bit signed int
uuid	<code>uuid</code>	A UUID (of any version)
varchar	<code>string</code>	UTF8 encoded string
varint	<code>integer</code>	Arbitrary-precision integer

- Collection types
 - `set<int>`
 - `list<text>`
 - `map<int, text>`
- User-Defined Types (UDTs)

```
CREATE TYPE address (
    street text,
    city text,
    state text,
);
```
- Tuple types
 - `tuple <uuid, text, decimal>`
- Custom types
 - Your own implementation as a Java class (not recommended)



More Types



Homework

Working with Data Types

[https://github.com/datastaxdevs/
workshop-cassandra-data-modeling](https://github.com/datastaxdevs/workshop-cassandra-data-modeling)

01



Data Storage Overview

Data Distribution & Organization

02

Key Definition

Primary, Partition, Clustering

03

CQL Data Types

Brief Review

04

Data Modeling Process

Methodology, Notation, Tool

05

Optimization Techniques

Five important optimizations

06

What's next?

Homework, Next Session



Agenda



Normalization vs Denormalization

"Database normalization is the process of structuring a relational database in accordance with a series of so-called normal forms in order to reduce data redundancy and improve data integrity. It was first proposed by Edgar F. Codd as part of his relational model."

PROS: Simple write, Data Integrity
CONS: Slow read, Complex Queries



Normalization

Employees

userId	deptId	firstName	lastName
1	1	Edgar	Codd
2	1	Raymond	Boyce

Departments

departmentId	department
1	Engineering
2	Math



"Denormalization is a strategy used on a database to increase performance. In computing, denormalization is the process of trying to improve the read performance of a database, at the expense of losing some write performance, by adding redundant copies of data"

PROS: Quick Read, Simple Queries

CONS: Multiple Writes, Manual Integrity

Employees

userId	firstName	lastName	department
1	Edgar	Codd	Engineering
2	Raymond	Boyce	Engineering
3	Sage	Lahja	Math
4	Juniper	Jones	Botany

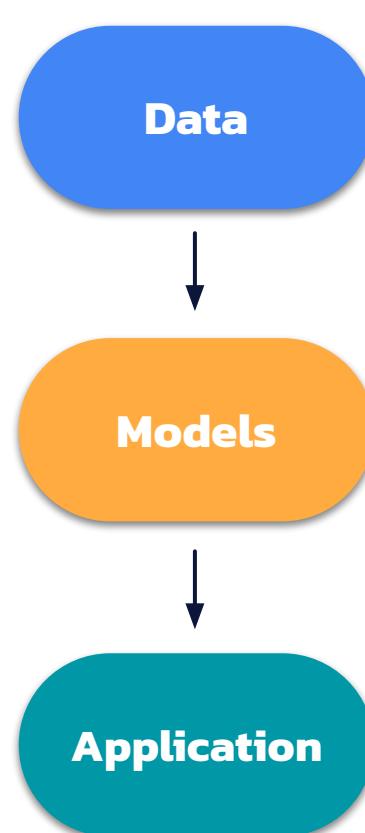
departmentId	department
1	Engineering
2	Math

Denormalization



Modelling Approach

1. Analyze raw data
2. Identify entities, their properties and relations
3. Design tables, using **normalization** and foreign keys.
4. Use JOIN when doing queries to join normalized data from multiple tables



A large cyan cylinder is positioned above two tables. A line connects the cylinder to the first table, and another line connects the cylinder to the second table.

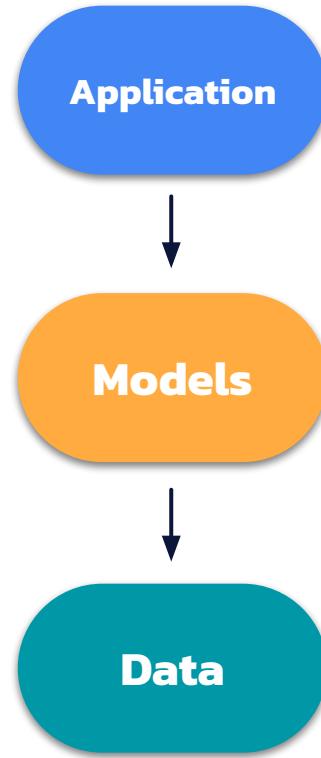
Employees		
userId	firstName	lastName
1	Edgar	Codd
2	Raymond	Boyce

Departments	
departmentId	department
1	Engineering
2	Math

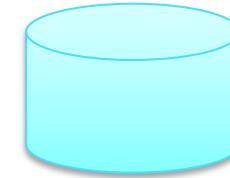


Relational Data Modelling

1. Analyze user behaviour
(customer first!)
2. Identify workflows, their dependencies and needs
3. Define Queries to fulfill these workflows
4. Knowing the queries, design tables, using **denormalization**.
5. Use BATCH when inserting or updating denormalized data of multiple tables



Employees			
userId	firstName	lastName	department
1	Edgar	Codd	Engineering
2	Raymond	Boyce	Math
3	Sage	Lahja	Math
4	Juniper	Jones	Botany



NoSQL Data Modelling

Data Modelling Methodology

by Dr. Artem Chebotko

Modelling process
Step by Step



- Collection and analysis of **data requirements**
- Identification of participating **entities and relationships**
- Identification of **data access patterns**
- A particular way of **organizing and structuring data**
- Design and specification of a **database schema**
- Schema **optimization** and data **indexing** techniques



Data Quality: completeness consistency accuracy
Data Access: queryability efficiency scalability



What is Data Modelling ?

Modeling principle 1: "Know your data"

- Key and cardinality constraints are fundamental to schema design

Modeling principle 2: "Know your queries"

- Queries drive schema design

Modeling principle 3: "Nest data"

- Data nesting is the main data modeling technique

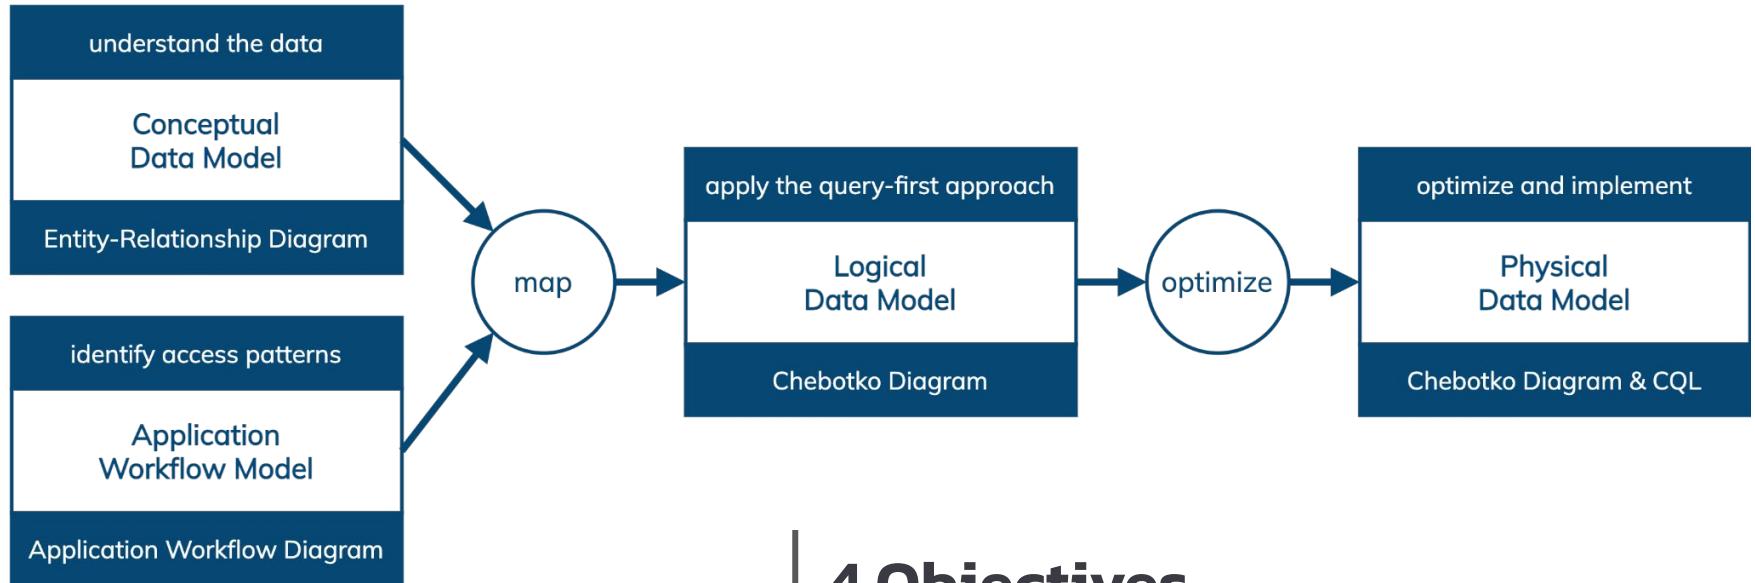
Modeling principle 4: "Duplicate data"

- Better to duplicate than to join



Cassandra Data Modelling Principles





**4 Objectives,
4 Models
2 Transitions**



4/4/2 Like at soccer (= *real* football)

understand the data

Conceptual Data Model

Entity-Relationship Diagram

Analyze the Domain

identify access patterns

Application Workflow Model

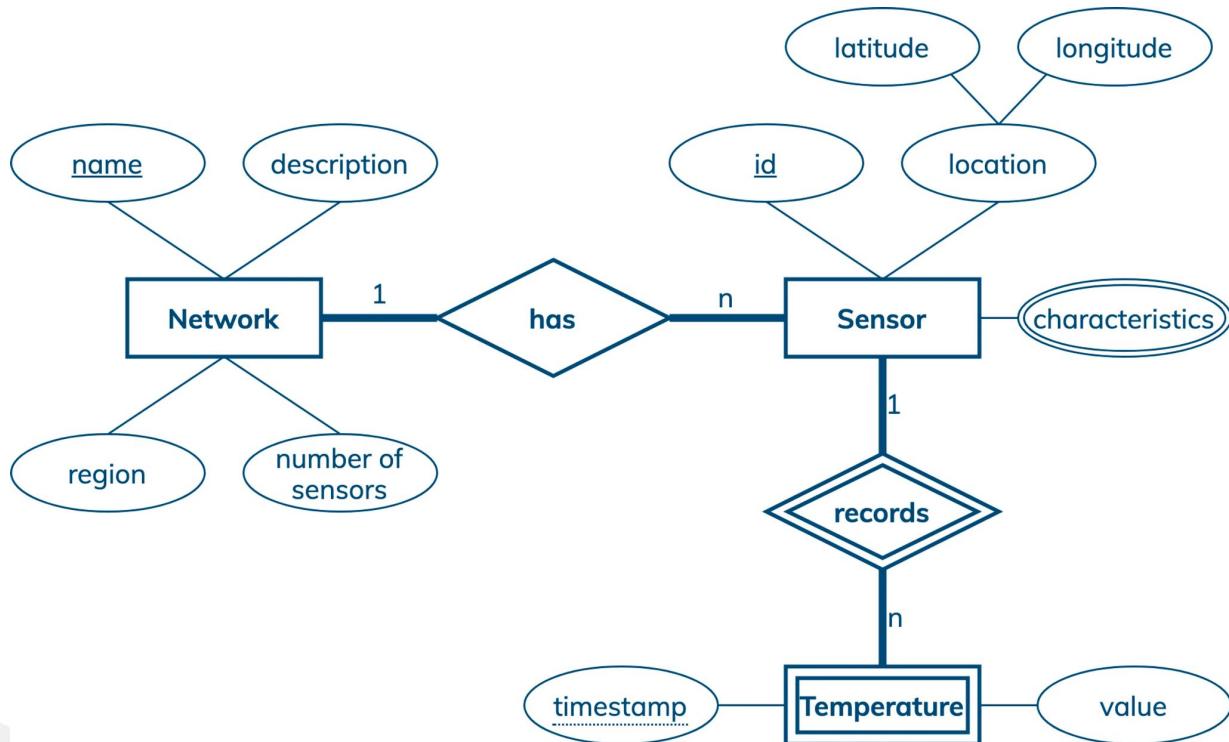
Application Workflow Diagram

Analyze Customer Workflows

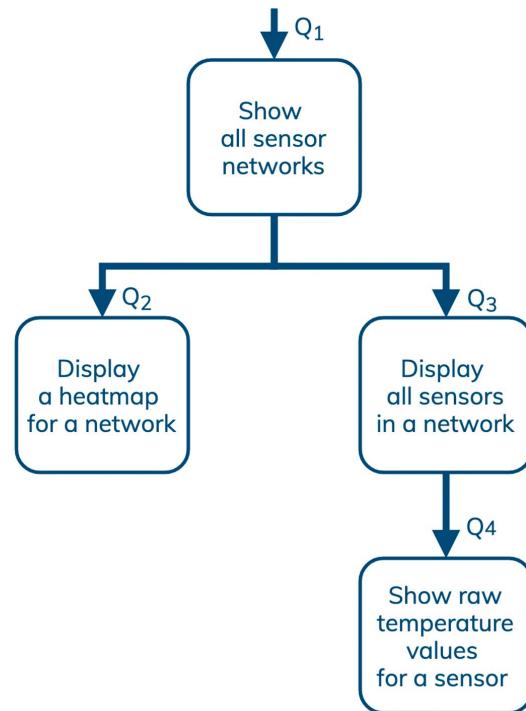


Data Modelling Methodology: Step I





Entity-Relationship Diagram



Data access patterns

Q1: Find information about all networks; order by name (asc)

Q2: Find hourly average temperatures for every sensor in a specified network for a given date range; order by date (desc) and hour (desc)

Q3: Find information about all sensors in a specified network

Q4: Find raw measurements for a particular sensor on a specified date; order by timestamp (desc)

Application Workflow Diagram

apply the query-first approach

Logical
Data Model

Chebotko Diagram

Design queries and build
tables based on the queries



Data Modelling Methodology: Step II

Mapping rule 1: "Entities and relationships"

- Entity and relationship types map to tables

Based on
a conceptual
data model

Mapping rule 2: "Equality search attributes"

- Equality search attributes map to the beginning columns of a primary key

Based on
a query

Mapping rule 3: "Inequality search attributes"

- Inequality search attributes map to clustering columns

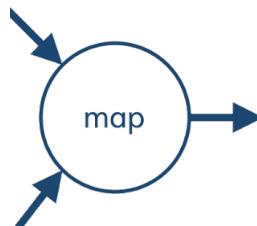
Mapping rule 4: "Ordering attributes"

- Ordering attributes map to clustering columns

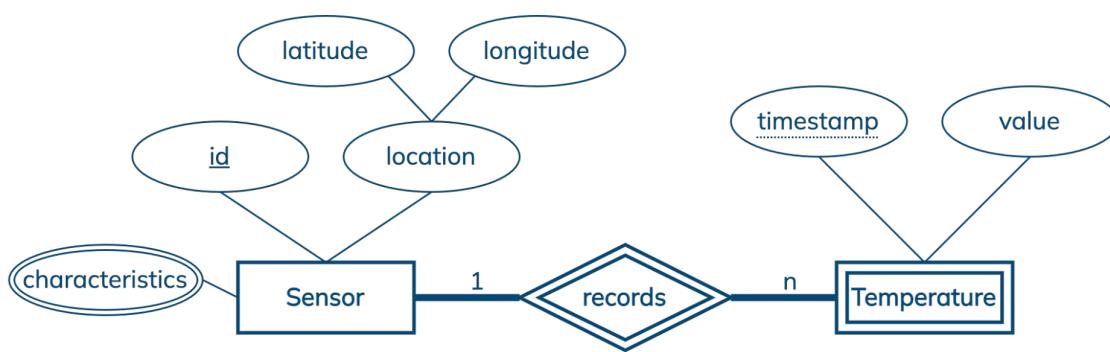
Based on
a conceptual
data model

Mapping rule 5: "Key attributes"

- Key attributes map to primary key columns



Mapping Rules

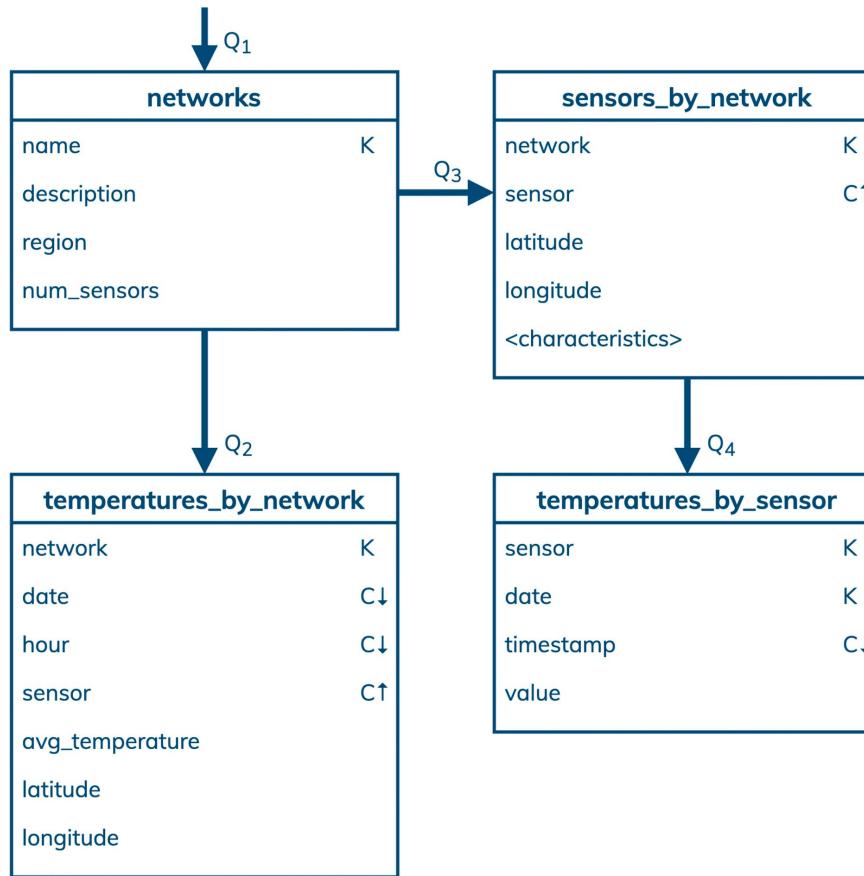


Q5: Find raw measurements for a given location and a date/time range; order by timestamp (desc)

The table shows the results of applying five mapping rules (MR1 to MR5) to the data. The columns represent different mapping configurations, and the rows represent the fields being mapped.

	MR ₁	MR ₂	MR ₃	MR ₄	MR ₅
temps_by_sensor	latitude longitude timestamp value	latitude longitude timestamp value	latitude longitude timestamp value	latitude longitude timestamp value	latitude longitude timestamp value
		K	K	C↑	K
		K	K	C↓	K
					C↑

Example: Applying Mapping Rules



Data access patterns

- Q1: Find information about all networks; order by name (asc)
- Q2: Find hourly average temperatures for every sensor in a specified network for a given date range; order by date (desc) and hour (desc)
- Q3: Find information about all sensors in a specified network
- Q4: Find raw measurements for a particular sensor on a specified date; order by timestamp (desc)



Chebotko Diagram



optimize and implement

Physical Data Model

Chebotko Diagram & CQL

Optimize and Create



Data Modelling Methodology: Step III

- Partition size limits and splitting large partitions
- Data duplication and batches
- Indexes and materialized views*
- Concurrent data access and lightweight transactions*
- Dealing with tombstones*

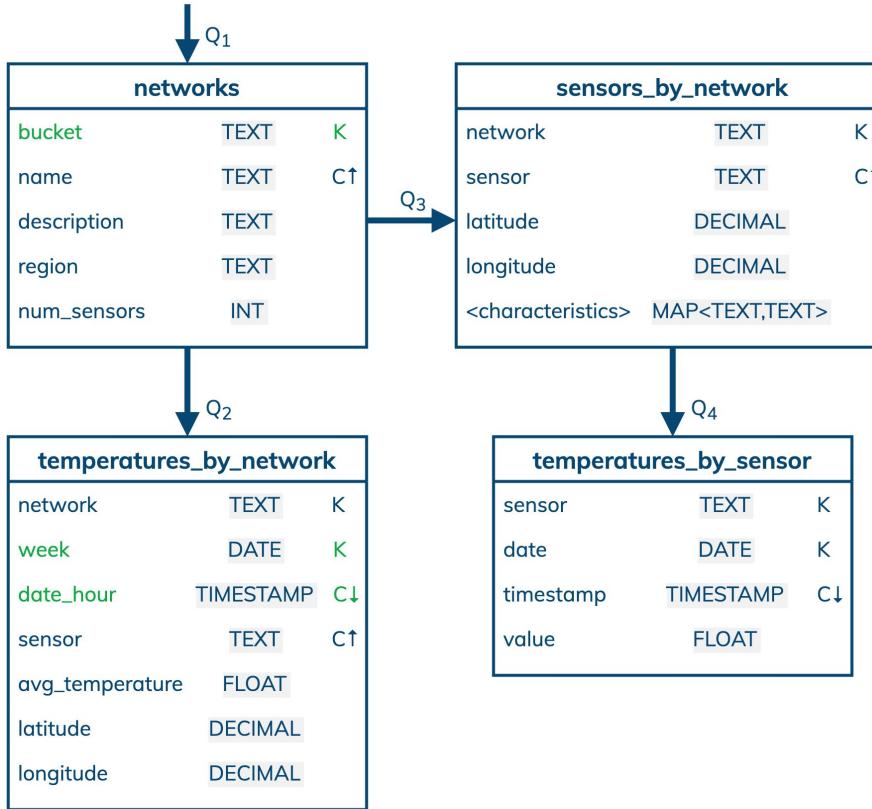


* Explained in details at our free DS220 course at academy.datastax.com



Optimization Techniques





Physical Data Model: Chebotko Diagram

Q1

```
CREATE TABLE networks (
    bucket TEXT,
    name TEXT,
    description TEXT,
    region TEXT,
    num_sensors INT,
    PRIMARY KEY ((bucket),name)
);
```

Q3

```
CREATE TABLE sensors_by_network (
    network TEXT,
    sensor TEXT,
    latitude DECIMAL,
    longitude DECIMAL,
    characteristics MAP<TEXT,TEXT>,
    PRIMARY KEY ((network),sensor)
);
```

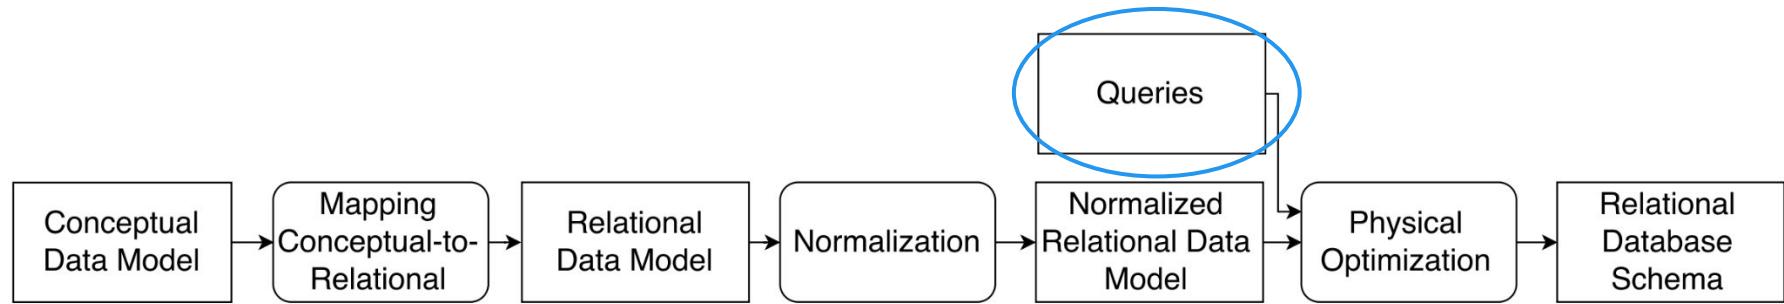
```
CREATE TABLE temperatures_by_network (
    network TEXT,
    week DATE,
    date_hour TIMESTAMP,
    sensor TEXT,
    avg_temperature FLOAT,
    latitude DECIMAL,
    longitude DECIMAL,
    PRIMARY KEY ((network,week),date_hour,sensor)
) WITH CLUSTERING ORDER BY (date_hour DESC, sensor ASC);
```

Q2

Q4

```
CREATE TABLE temperatures_by_sensor (
    sensor TEXT,
    date DATE,
    timestamp TIMESTAMP,
    value FLOAT,
    PRIMARY KEY ((sensor,date),timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```

Cassandra Query Language



(a) Relational Data Modeling



(b) Cassandra Data Modeling

Relational vs Cassandra Modeling Process



Demo

KDM

[https://github.com/datastaxdevs/
workshop-cassandra-data-modeling](https://github.com/datastaxdevs/workshop-cassandra-data-modeling)





Lab 3 (Optional)

Sensor Data Modeling

[https://github.com/datastaxdevs/
workshop-cassandra-data-modeling](https://github.com/datastaxdevs/workshop-cassandra-data-modeling)

Data Modeling Principles

- Know your data
- Know your queries
- Nest data
- Duplicate data

The Most Efficient Access Pattern

A query is satisfied by accessing one partition

Primary Key

- Data uniqueness
- Data distribution
- Data queryability

The Cassandra Data Modeling Methodology

It is not more complex than the relational data modeling methodology



Key Takeaways

01



Data Storage Overview
Data Distribution & Organization

02

Key Definition
Primary, Partition, Clustering

03

CQL Data Types
Brief Review

04

Data Modeling Process
Methodology, Notation, Tool

05

Optimization Techniques
Five important optimizations

06

What's next?
Homework, Next Session



Agenda

Partition size limits and splitting large partitions



- **Theoretical**
 - 2 billion values
 - Node disk size
- **Practical**

Cassandra 2	Cassandra 3 and 4
100,000 values	10 x 100,000 values
100 MB	10 x 100 MB

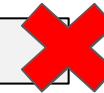


Partition Size Limits



- **Solution: introduce an additional column to a partition key**
 - Use an existing or derived column – convenience
 - Use an artificial “bucket” column – more control
- **Cons: supported access patterns may change**

```
PRIMARY KEY ((sensor), timestamp);
```



```
PRIMARY KEY ((sensor, month), timestamp);
```



Data duplication and batches



- **Types of data duplication**
 - Duplication across tables
 - Duplication across partitions of the same table
 - Duplication across rows of the same partition
- **Duplication factor: constant ✓ vs non-constant ✗**



Data Duplication



- **Logged, single-partition batches**
 - Efficient and atomic
 - Use whenever possible
- **Logged, multi-partition batches**
 - Somewhat efficient and pseudo-atomic
 - Use selectively
- **Unlogged and counter batches**
 - Do not use

```
BEGIN BATCH  
  INSERT INTO ...;  
  INSERT INTO ...;  
  ...  
  APPLY BATCH;  
BEGIN BATCH  
  UPDATE ...;  
  UPDATE ...;  
  ...  
  APPLY BATCH;
```



Keeping Duplicated Data Consistent with Batches



- **Experimental feature**
- **Not always applicable**
- **May become inconsistent**
- **10% write penalty to a base table**

```
CREATE MATERIALIZED VIEW  
networks_by_region AS  
    SELECT * FROM networks  
    WHERE region IS NOT NULL  
        AND name IS NOT NULL  
PRIMARY KEY ((region), name);
```



Keeping Duplicated Data Consistent with MVs

Indexes and materialized views (+ allow filtering)

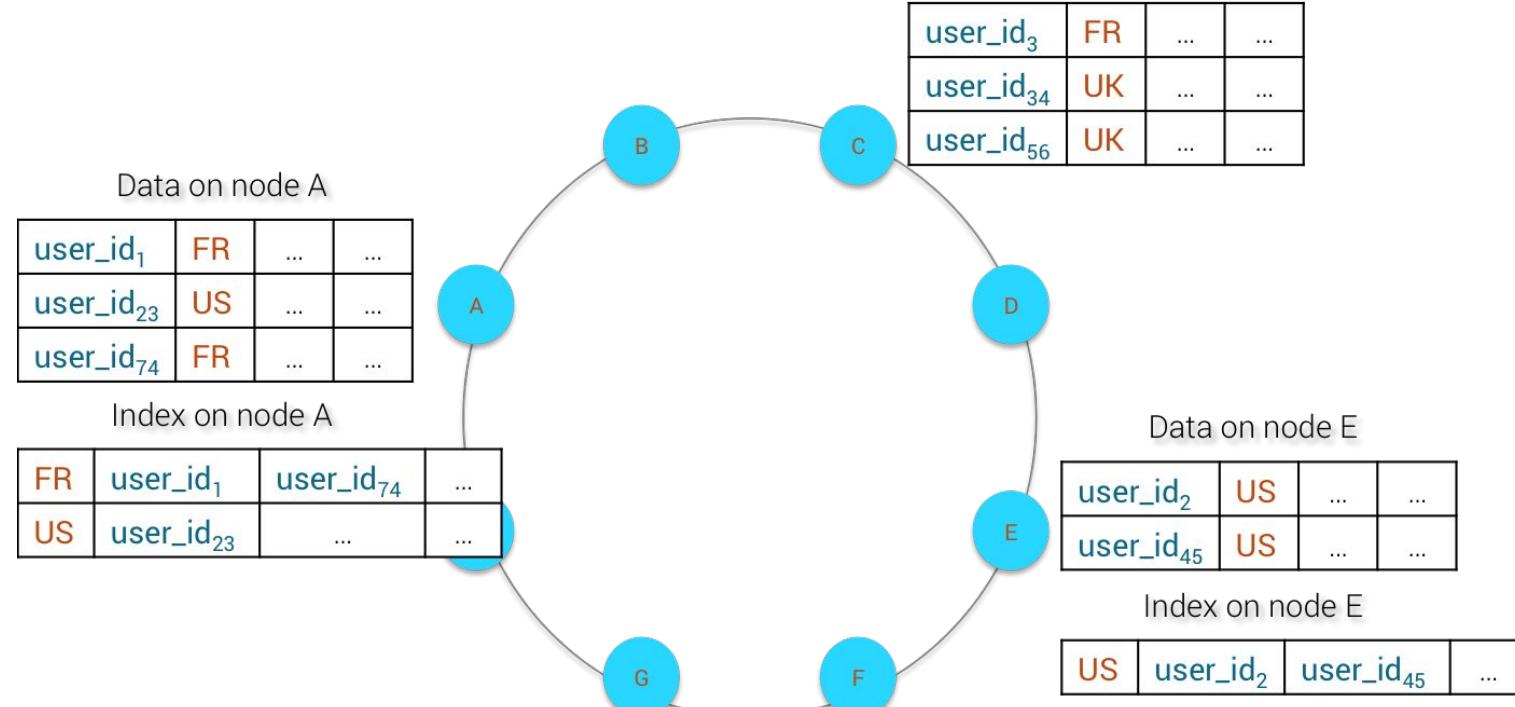


- **Distributed indexes: highly scalable**
 - **Tables ← the best solution available**
 - **Materialized views (experimental)**
- **Local indexes: not highly scalable**
 - **Secondary indexes (2i)**
 - **SSTable-attached secondary indexes (SASI)**
 - **Storage-Attached Indexes (SAI)**



Indexing Mechanisms





doanduyhai.com/blog/?p=13191

How Secondary Indexes Work

- **The best index available**
 - Faster writes
 - Less disk space
 - Multiple indexes can be used in one query
 - Reads can be expensive (similar limitations to 2is and SASIs)
- **Convenience, not performance**

```
CREATE CUSTOM INDEX users_by_name_sai  
ON users (name)  
USING 'org.apache.cassandra.index.sai.StorageAttachedIndex' ;
```



Storage-Attached Indexes (SAI)



- **Real-time transactional query**: retrieving rows from a large multi-row partition, when both a **partition key** and an indexed column are used in a query.
- **Expensive analytical query**: retrieving a large number of rows from potentially many partitions, when only an indexed **low-cardinality column** is used in a query.



When to use secondary index (1 of 2)

- **Small clusters**: 3–10 nodes with RF=3
- **Infrequent queries**: may not affect the throughput much
- **Queries with LIMIT that can be quickly satisfied**: may not affect the throughput much
- **Convenience over performance**: simpler data modeling while potentially sacrificing throughput and query response time



When to use secondary index (2 of 2)



- **Queries on higher-cardinality columns**
 - Similar advantages as those of regular tables
 - Convenience of automatic view maintenance
 - Reads are as fast as for regular tables
- **Important limitations**
 - Restrictions on how PRIMARY KEY is constructed
 - 10% slower writes to the base table
 - Base-view inconsistencies
 - Use LOCAL_QUORUM and higher for base table writes
 - Disabled in Astra DB, experimental



When to Use a Materialized View

- Generally, using **ALLOW FILTERING** is not recommended
- But scanning may be cheaper than indexing in some cases
 - Scanning within a small partition
 - Scanning within a large partition based on a low-cardinality column (most rows from the partition will be returned)
 - Scanning a table based on a low-cardinality column (most rows from the table will be returned) for analytical purposes



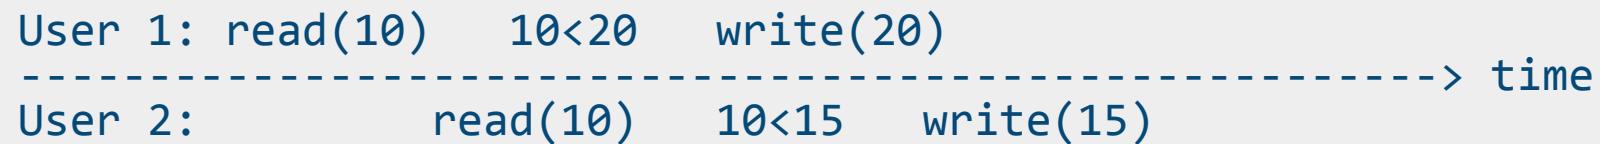
Allow Filtering



Concurrent data access and lightweight transactions



- Multiple transactions reading and writing the same piece of data concurrently, possibly leaving the data in an inconsistent-with-reality state
- Examples: registering a new username, counting votes, updating account balance, updating an order status, updating the highest bid



Race conditions



- **Linearizable consistency ensures that concurrent transactions produce the same result as if they execute in a sequence, one after another**
- **Lightweight transactions (LWTs) in Cassandra and Astra DB**

User 1: read(10) 10<20 write(20)

-----> time

User 2: read(20) 20>15

User 1: read(15) 15<20 write(20)

-----> time

User 2: read(10) 10<15 write(15)



Linearizable Consistency



- **Guarantee linearizable consistency**
- **Use with race conditions and low data contention**
- **Require 4 coordinator-replica round trips**
- **May become a bottleneck with high data contention**

```
INSERT INTO ... VALUES ...
IF NOT EXISTS;

UPDATE ... SET ... WHERE ...
IF EXISTS | IF predicate [ AND ... ];

DELETE ... FROM ... WHERE ...
IF EXISTS | IF predicate [ AND ... ];
```



Lightweight transactions (LWTs)

- **Change a data model to eliminate race conditions**
 - Isolate computation by isolating data
 - Record each transaction separately and aggregate data later
- **Use a queue to control data contention**
 - Astra Streaming



Dealing with High Data Contention



Tombstones



- **Written markers for deleted data**
 - **DELETEs**
 - **Data with TTL**
 - **UPDATEs and INSERTs with NULL values**
 - **“Overwritten” collections**



Tombstones



- Tombstone = additional data to store and read
- Query performance degrades, heap memory pressure increases
- **tombstone_warn_threshold**
 - Warning when more than 1,000 tombstones are scanned by a query
- **tombstone_failure_threshold**
 - Aborted query when more than 100,000 tombstones are scanned



1. Decrease the value of **gc_grace_seconds** (default is 864000 or 10 days)
 - Deleted data and tombstones can be purged during compaction after **gc_grace_seconds**
2. Run compaction more frequently
 - **nodetool compact keyspace tablename**



Tombstones Issue 1: Large Partitions

1. A replica node is unresponsive and receives no tombstone
2. Other replica nodes receive tombstones
3. The tombstones get purged after `gc_grace_seconds` and compaction
4. The unresponsive replica comes back online and resurrects data that was previously marked as deleted



- Run repairs within the `gc_grace_seconds` and on a regular basis
- `nodetool repair`
- Do not let a node to rejoin the cluster after the `gc_grace_seconds`



Tombstones Issue 2: Zombie or Resurrected Data

- **Design data models that avoid frequent deletes**
- **Truncating/dropping a table is better than deleting each partition in the table**
- **Deleting a partition is better than deleting each row in the partition**
- **Deleting a row is better than deleting each value in the row**



Other Tombstones Considerations



Partition Size

100K and 100MB

Bucketing

Data duplication and batches

Avoid non-constant duplication factors

Single-partition batches are efficient

Indexing

Materialized views are scalable but have other issues

Secondary indexes are good for searching within large partitions

LWTs

Linearizable consistency

Race conditions and low data contention

Design data models that eliminate race conditions

Tombstones

Design data models that avoid frequent deletions

Compact and repair to purge tombstones and avoid resurrected data



01



Data Storage Overview
Data Distribution & Organization

02

Key Definition
Primary, Partition, Clustering

03

CQL Data Types
Brief Review

04

Data Modeling Process
Methodology, Notation, Tool

05

Optimization Techniques
Five important optimizations

06

What's next?
Homework, Next Session



Agenda

datastax.com/learn/data-modeling-by-example

killercoda.com/datastaxdevs/course/cassandra-data-modeling



A grid of nine cards, each representing a different data modeling scenario:

- Sensor Data Modeling: Learn how to create a data model for temperature monitoring sensor networks.
- Messaging Data Modeling: Learn how to create a data model for an email system.
- Digital Library Data Modeling: Learn how to create a data model for a digital music library.
- Investment Portfolio Data Modeling: Learn how to create a data model for investment accounts or portfolios.
- Time Series Data Modeling: Learn how to create a data model for time series data.
- Shopping Cart Data Modeling: Learn how to create a data model for online shopping carts.
- Order Management Data Modeling: Learn how to create a data model for an order management system.



HomeWork

- medium.com/building-the-open-data-stack/data-modeling-in-cassandra-and-datastax-astra-db-3f89b9c133c9
- devm.io/iot/apache-cassandra-iot-174970



Blog Posts





Cassandra Data Modelling Homework



Welcome and thank you! Here you can submit your homework for the datastax developers "Data Modelling with Cassandra" workshop. In case of any questions please contact organizers at <https://dtsx.io/aleks> or just send an email to aleksandr.volochnev@datastax.com

- Workshop materials: <https://github.com/datastaxdevs/workshop-intro-to-cassandra>
- Discord chat: <https://dtsx.io/discord>

Email *

Valid email

This form is collecting emails. [Change settings](#)



Join our 19k Discord Community

DataStax Developers



!discord

dtsx.io/discord

The screenshot shows the DataStax Developers Discord server interface. The left sidebar lists various channels: Événements, # moderator-only, # WELCOME, start-here, code-of-conduct, # introductions, upcoming-events, useful-resources, # memes, # your-ideas, @ the-stage, # WORKSHOPS, # workshop-chat, # workshop-feedback, # workshop-materials, # upcoming-workshops, # ASTRADB, # getting-started, # astra-apis, # astra-development, # sample-applications, and # APACHE CASSANDRA. The main window is the # workshop-chat channel, which contains a message from a user named RIGGITYREKT asking about mixed version testing for a class and one node that is 5.0.15. Another message from Erick Ramirez follows, stating that mixed versions are not supported. The right sidebar shows a list of presenters and helpers, along with a list of members currently online.

Discord Community

Thank You!

